

CNN – Classification Dogs vs Cats

Objectif : entraîner un CNN pour une classification binaire et analyser l'impact du nombre de convolutions.

```
In [24]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
import os
from glob import glob
```

```
In [2]: IMG_SIZE = (64, 64)
BATCH_SIZE = 32
EPOCHS = 20
```

```
In [4]: train_gen = ImageDataGenerator(
    rescale=1./255,
    zoom_range=0.2,
    shear_range=0.2,
    horizontal_flip=True
)

test_gen = ImageDataGenerator(rescale=1./255)

train_set = train_gen.flow_from_directory(
    "../dataset/training_set",
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="binary"
)

test_set = test_gen.flow_from_directory(
    "../dataset/test_set",
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="binary"
)
```

Found 8000 images belonging to 2 classes.
 Found 2000 images belonging to 2 classes.

```
In [9]: def build_cnn(nb_conv):
    model = Sequential()

    # 1ère convolution
    model.add(Conv2D(32, (3,3), activation='relu', input_shape=(64,64,3)))
    model.add(MaxPooling2D((2,2)))

    # 2ème convolution
    model.add(Conv2D(32, (3,3), activation='relu'))
    model.add(MaxPooling2D((2,2)))
```

```
# 3ème convolution OPTIONNELLE
if nb_conv == 3:
    model.add(Conv2D(64, (3,3), activation='relu'))
    model.add(MaxPooling2D((2,2)))

    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)

return model
```

In [20]:

```
def get_callbacks(nb_conv):
    ckpt_dir = f"checkpoints_cnn_{nb_conv}"
    os.makedirs(ckpt_dir, exist_ok=True)

    checkpoint = ModelCheckpoint(
        filepath=os.path.join(ckpt_dir, "weights_epoch_{epoch:02d}.h5"),
        monitor="val_loss",
        save_weights_only=True,
        save_best_only=False,
        verbose=1
    )

    early_stop = EarlyStopping(
        monitor="val_loss",
        patience=3,
        restore_best_weights=True,
        verbose=1
    )

    return checkpoint, early_stop, ckpt_dir
```

In [21]:

```
model_2 = build_cnn(nb_conv=2)

checkpoint, early_stop, ckpt_dir = get_callbacks(2)

history_2 = model_2.fit(
    train_set,
    steps_per_epoch=train_set.samples // BATCH_SIZE,
    epochs=20,
    validation_data=test_set,
    validation_steps=test_set.samples // BATCH_SIZE,
    callbacks=[checkpoint, early_stop],
    verbose=1
)
```

```
Epoch 1/20
250/250 [=====] - ETA: 0s - loss: 0.6654 - accuracy: 0.5949
Epoch 1: saving model to checkpoints_cnn_2\weights_epoch_01.h5
250/250 [=====] - 29s 115ms/step - loss: 0.6654 - accuracy: 0.5949 - val_loss: 0.6091 - val_accuracy: 0.6628
Epoch 2/20
250/250 [=====] - ETA: 0s - loss: 0.6124 - accuracy: 0.626
Epoch 2: saving model to checkpoints_cnn_2\weights_epoch_02.h5
250/250 [=====] - 33s 132ms/step - loss: 0.6124 - accuracy: 0.6626 - val_loss: 0.5721 - val_accuracy: 0.7107
Epoch 3/20
250/250 [=====] - ETA: 0s - loss: 0.5716 - accuracy: 0.6991
Epoch 3: saving model to checkpoints_cnn_2\weights_epoch_03.h5
250/250 [=====] - 45s 178ms/step - loss: 0.5716 - accuracy: 0.6991 - val_loss: 0.5426 - val_accuracy: 0.7253
Epoch 4/20
250/250 [=====] - ETA: 0s - loss: 0.5367 - accuracy: 0.7258
Epoch 4: saving model to checkpoints_cnn_2\weights_epoch_04.h5
250/250 [=====] - 53s 212ms/step - loss: 0.5367 - accuracy: 0.7258 - val_loss: 0.5347 - val_accuracy: 0.7324
Epoch 5/20
250/250 [=====] - ETA: 0s - loss: 0.5122 - accuracy: 0.7455
Epoch 5: saving model to checkpoints_cnn_2\weights_epoch_05.h5
250/250 [=====] - 31s 123ms/step - loss: 0.5122 - accuracy: 0.7455 - val_loss: 0.4911 - val_accuracy: 0.7676
Epoch 6/20
250/250 [=====] - ETA: 0s - loss: 0.4926 - accuracy: 0.7565
Epoch 6: saving model to checkpoints_cnn_2\weights_epoch_06.h5
250/250 [=====] - 31s 125ms/step - loss: 0.4926 - accuracy: 0.7565 - val_loss: 0.4779 - val_accuracy: 0.7732
Epoch 7/20
250/250 [=====] - ETA: 0s - loss: 0.4766 - accuracy: 0.7625
Epoch 7: saving model to checkpoints_cnn_2\weights_epoch_07.h5
250/250 [=====] - 29s 116ms/step - loss: 0.4766 - accuracy: 0.7625 - val_loss: 0.4770 - val_accuracy: 0.7818
Epoch 8/20
250/250 [=====] - ETA: 0s - loss: 0.4562 - accuracy: 0.7780
Epoch 8: saving model to checkpoints_cnn_2\weights_epoch_08.h5
250/250 [=====] - 29s 114ms/step - loss: 0.4562 - accuracy: 0.7780 - val_loss: 0.5107 - val_accuracy: 0.7530
Epoch 9/20
250/250 [=====] - ETA: 0s - loss: 0.4455 - accuracy: 0.7905
Epoch 9: saving model to checkpoints_cnn_2\weights_epoch_09.h5
250/250 [=====] - 27s 109ms/step - loss: 0.4455 - accuracy: 0.7905 - val_loss: 0.4998 - val_accuracy: 0.7651
Epoch 10/20
250/250 [=====] - ETA: 0s - loss: 0.4305 - accuracy: 0.7971
Epoch 10: saving model to checkpoints_cnn_2\weights_epoch_10.h5
250/250 [=====] - 32s 127ms/step - loss: 0.4305 - accuracy: 0.7971 - val_loss: 0.4558 - val_accuracy: 0.7878
```

```
Epoch 11/20
250/250 [=====] - ETA: 0s - loss: 0.4223 - accuracy: 0.8058
Epoch 11: saving model to checkpoints_cnn_2\weights_epoch_11.h5
250/250 [=====] - 35s 140ms/step - loss: 0.4223 - accuracy: 0.8058 - val_loss: 0.4966 - val_accuracy: 0.7722
Epoch 12/20
250/250 [=====] - ETA: 0s - loss: 0.4076 - accuracy: 0.8123
Epoch 12: saving model to checkpoints_cnn_2\weights_epoch_12.h5
250/250 [=====] - 47s 188ms/step - loss: 0.4076 - accuracy: 0.8123 - val_loss: 0.4697 - val_accuracy: 0.7883
Epoch 13/20
250/250 [=====] - ETA: 0s - loss: 0.4013 - accuracy: 0.8110
Epoch 13: saving model to checkpoints_cnn_2\weights_epoch_13.h5
Restoring model weights from the end of the best epoch: 10.
250/250 [=====] - 35s 140ms/step - loss: 0.4013 - accuracy: 0.8110 - val_loss: 0.5542 - val_accuracy: 0.7450
Epoch 13: early stopping
```

```
In [22]: model_3 = build_cnn(nb_conv=3)

checkpoint, early_stop, ckpt_dir = get_callbacks(3)

history_3 = model_3.fit(
    train_set,
    steps_per_epoch=train_set.samples // BATCH_SIZE,
    epochs=20,
    validation_data=test_set,
    validation_steps=test_set.samples // BATCH_SIZE,
    callbacks=[checkpoint, early_stop],
    verbose=1
)
```

```
Epoch 1/20
250/250 [=====] - ETA: 0s - loss: 0.6786 - accuracy: 0.5734
Epoch 1: saving model to checkpoints_cnn_3\weights_epoch_01.h5
250/250 [=====] - 31s 120ms/step - loss: 0.6786 - accuracy: 0.5734 - val_loss: 0.6485 - val_accuracy: 0.6295
Epoch 2/20
250/250 [=====] - ETA: 0s - loss: 0.6380 - accuracy: 0.6424
Epoch 2: saving model to checkpoints_cnn_3\weights_epoch_02.h5
250/250 [=====] - 30s 121ms/step - loss: 0.6380 - accuracy: 0.6424 - val_loss: 0.6041 - val_accuracy: 0.6764
Epoch 3/20
250/250 [=====] - ETA: 0s - loss: 0.5917 - accuracy: 0.6821
Epoch 3: saving model to checkpoints_cnn_3\weights_epoch_03.h5
250/250 [=====] - 29s 116ms/step - loss: 0.5917 - accuracy: 0.6821 - val_loss: 0.5744 - val_accuracy: 0.7001
Epoch 4/20
250/250 [=====] - ETA: 0s - loss: 0.5566 - accuracy: 0.7109
Epoch 4: saving model to checkpoints_cnn_3\weights_epoch_04.h5
250/250 [=====] - 29s 117ms/step - loss: 0.5566 - accuracy: 0.7109 - val_loss: 0.5384 - val_accuracy: 0.7349
Epoch 5/20
250/250 [=====] - ETA: 0s - loss: 0.5349 - accuracy: 0.7243
Epoch 5: saving model to checkpoints_cnn_3\weights_epoch_05.h5
250/250 [=====] - 29s 115ms/step - loss: 0.5349 - accuracy: 0.7243 - val_loss: 0.5740 - val_accuracy: 0.7001
Epoch 6/20
250/250 [=====] - ETA: 0s - loss: 0.5120 - accuracy: 0.7475
Epoch 6: saving model to checkpoints_cnn_3\weights_epoch_06.h5
250/250 [=====] - 28s 113ms/step - loss: 0.5120 - accuracy: 0.7475 - val_loss: 0.4977 - val_accuracy: 0.7661
Epoch 7/20
250/250 [=====] - ETA: 0s - loss: 0.4908 - accuracy: 0.7594
Epoch 7: saving model to checkpoints_cnn_3\weights_epoch_07.h5
250/250 [=====] - 29s 115ms/step - loss: 0.4908 - accuracy: 0.7594 - val_loss: 0.4848 - val_accuracy: 0.7712
Epoch 8/20
250/250 [=====] - ETA: 0s - loss: 0.4728 - accuracy: 0.7713
Epoch 8: saving model to checkpoints_cnn_3\weights_epoch_08.h5
250/250 [=====] - 32s 129ms/step - loss: 0.4728 - accuracy: 0.7713 - val_loss: 0.4999 - val_accuracy: 0.7616
Epoch 9/20
250/250 [=====] - ETA: 0s - loss: 0.4554 - accuracy: 0.7879
Epoch 9: saving model to checkpoints_cnn_3\weights_epoch_09.h5
250/250 [=====] - 28s 113ms/step - loss: 0.4554 - accuracy: 0.7879 - val_loss: 0.4841 - val_accuracy: 0.7636
Epoch 10/20
250/250 [=====] - ETA: 0s - loss: 0.4319 - accuracy: 0.7956
Epoch 10: saving model to checkpoints_cnn_3\weights_epoch_10.h5
250/250 [=====] - 30s 120ms/step - loss: 0.4319 - accuracy: 0.7956 - val_loss: 0.4613 - val_accuracy: 0.7949
```

```

Epoch 11/20
250/250 [=====] - ETA: 0s - loss: 0.4254 - accuracy: 0.8030
Epoch 11: saving model to checkpoints_cnn_3\weights_epoch_11.h5
250/250 [=====] - 38s 152ms/step - loss: 0.4254 - accuracy: 0.8030 - val_loss: 0.4453 - val_accuracy: 0.7979
Epoch 12/20
250/250 [=====] - ETA: 0s - loss: 0.4064 - accuracy: 0.8120
Epoch 12: saving model to checkpoints_cnn_3\weights_epoch_12.h5
250/250 [=====] - 46s 183ms/step - loss: 0.4064 - accuracy: 0.8120 - val_loss: 0.4343 - val_accuracy: 0.8019
Epoch 13/20
250/250 [=====] - ETA: 0s - loss: 0.3835 - accuracy: 0.8282
Epoch 13: saving model to checkpoints_cnn_3\weights_epoch_13.h5
250/250 [=====] - 46s 185ms/step - loss: 0.3835 - accuracy: 0.8282 - val_loss: 0.4491 - val_accuracy: 0.7979
Epoch 14/20
250/250 [=====] - ETA: 0s - loss: 0.3783 - accuracy: 0.8270
Epoch 14: saving model to checkpoints_cnn_3\weights_epoch_14.h5
250/250 [=====] - 42s 169ms/step - loss: 0.3783 - accuracy: 0.8270 - val_loss: 0.4142 - val_accuracy: 0.8140
Epoch 15/20
250/250 [=====] - ETA: 0s - loss: 0.3592 - accuracy: 0.8396
Epoch 15: saving model to checkpoints_cnn_3\weights_epoch_15.h5
250/250 [=====] - 42s 168ms/step - loss: 0.3592 - accuracy: 0.8396 - val_loss: 0.4990 - val_accuracy: 0.7742
Epoch 16/20
250/250 [=====] - ETA: 0s - loss: 0.3500 - accuracy: 0.8445
Epoch 16: saving model to checkpoints_cnn_3\weights_epoch_16.h5
250/250 [=====] - 36s 143ms/step - loss: 0.3500 - accuracy: 0.8445 - val_loss: 0.4429 - val_accuracy: 0.8019
Epoch 17/20
250/250 [=====] - ETA: 0s - loss: 0.3325 - accuracy: 0.8547
Epoch 17: saving model to checkpoints_cnn_3\weights_epoch_17.h5
Restoring model weights from the end of the best epoch: 14.
250/250 [=====] - 33s 131ms/step - loss: 0.3325 - accuracy: 0.8547 - val_loss: 0.4498 - val_accuracy: 0.8044
Epoch 17: early stopping

```

```

In [23]: import matplotlib.pyplot as plt

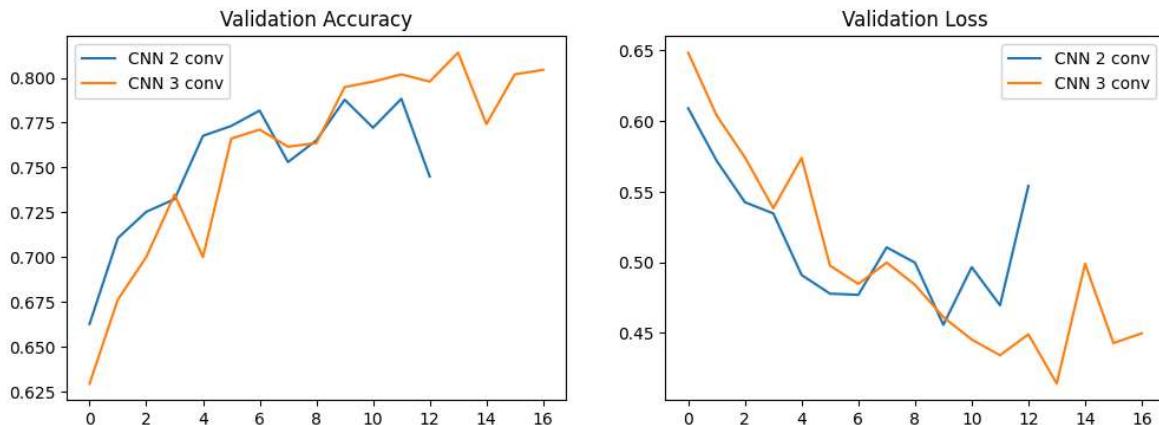
plt.figure(figsize=(12,4))

plt.subplot(1,2,1)
plt.plot(history_2.history["val_accuracy"], label="CNN 2 conv")
plt.plot(history_3.history["val_accuracy"], label="CNN 3 conv")
plt.title("Validation Accuracy")
plt.legend()

plt.subplot(1,2,2)
plt.plot(history_2.history["val_loss"], label="CNN 2 conv")
plt.plot(history_3.history["val_loss"], label="CNN 3 conv")
plt.title("Validation Loss")
plt.legend()

```

```
plt.show()
```



Comparaison CNN 2 convolutions vs CNN 3 convolutions (Dogs & Cats)

L'entraînement des deux architectures montre que l'ajout d'une troisième couche de convolution améliore légèrement les performances, tout en maintenant un modèle stable et bien régularisé grâce à l'early stopping.

CNN 2 convolutions :

- Précision maximale sur l'ensemble de validation : ~78,2 %
- Perte de validation minimale : ~0,456
- Early stopping déclenché à l'epoch 13, meilleur modèle restauré à l'epoch 10

CNN 3 convolutions :

- Précision maximale sur l'ensemble de validation : ~80,2 %
- Perte de validation minimale : ~0,434
- Early stopping déclenché à l'epoch 17, meilleur modèle restauré à l'epoch 14

Le graphique comparatif montre que CNN 3 convolutions converge légèrement plus rapidement et atteint une précision supérieure, ce qui confirme que la couche convolutionnelle supplémentaire permet au modèle de mieux capturer les caractéristiques des images, même pour une tâche binaire comme Dogs & Cats.

Donc l'ajout de la troisième convolution améliore la performance globale du CNN, sans provoquer d'overfitting significatif. Pour ce dataset, le CNN à 3 convolutions est donc légèrement préférable.