

CNN – Classification Intel Image Classification

Objectif : entraîner un CNN pour une classification binaire et analyser l'impact du nombre de convolutions.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
import os
from glob import glob
```

```
In [10]: IMG_SIZE = (150, 150)
BATCH_SIZE = 32
EPOCHS = 20
```

```
In [11]: train_datagen = ImageDataGenerator(
    rescale=1./255,
    zoom_range=0.2,
    shear_range=0.2,
    horizontal_flip=True
)

validation_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    "../dataset/seg_train",
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="categorical" # MULTICLASSE
)

validation_generator = validation_datagen.flow_from_directory(
    "../dataset/seg_test",
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode="categorical" # MULTICLASSE
)
```

Found 14034 images belonging to 6 classes.

Found 3000 images belonging to 6 classes.

```
In [12]: def build_cnn(nb_conv):
    model = Sequential()

    # 1ère convolution
    model.add(Conv2D(32, (3,3), activation='relu', input_shape=(150,150,3)))
    model.add(MaxPooling2D((2,2)))

    # 2ème convolution
    model.add(Conv2D(32, (3,3), activation='relu'))
    model.add(MaxPooling2D((2,2)))
```

```

# 3ème convolution OPTIONNELLE
if nb_conv == 3:
    model.add(Conv2D(64, (3,3), activation='relu'))
    model.add(MaxPooling2D((2,2)))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))           # régularisation pour éviter overfitting
model.add(Dense(6, activation='softmax')) # 6 classes pour Intel

model.compile(
    optimizer='adam',
    loss='categorical_crossentropy', # multiclass
    metrics=['accuracy']
)

return model

```

```

In [13]: def get_callbacks(nb_conv):
    ckpt_dir = f"checkpoints_cnn_{nb_conv}"
    os.makedirs(ckpt_dir, exist_ok=True)

    checkpoint = ModelCheckpoint(
        filepath=os.path.join(ckpt_dir, "weights_epoch_{epoch:02d}.h5"),
        monitor="val_accuracy",
        save_weights_only=True,
        save_best_only=True,
        verbose=1
    )

    early_stop = EarlyStopping(
        monitor="val_accuracy",
        patience=5,
        restore_best_weights=True,
        verbose=1
    )

    return checkpoint, early_stop, ckpt_dir

```

```

In [14]: model_2 = build_cnn(nb_conv=2)

checkpoint, early_stop, ckpt_dir = get_callbacks(2)

history_2 = model_2.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // BATCH_SIZE,
    epochs=20,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // BATCH_SIZE,
    callbacks=[checkpoint, early_stop],
    verbose=1
)

```

```
Epoch 1/20
438/438 [=====] - ETA: 0s - loss: 1.1739 - accuracy: 0.5390
Epoch 1: val_accuracy improved from -inf to 0.67641, saving model to checkpoints_cnn_2\weights_epoch_01.h5
438/438 [=====] - 292s 666ms/step - loss: 1.1739 - accuracy: 0.5390 - val_loss: 0.8306 - val_accuracy: 0.6764
Epoch 2/20
438/438 [=====] - ETA: 0s - loss: 0.9267 - accuracy: 0.6525
Epoch 2: val_accuracy did not improve from 0.67641
438/438 [=====] - 174s 398ms/step - loss: 0.9267 - accuracy: 0.6525 - val_loss: 0.9686 - val_accuracy: 0.6378
Epoch 3/20
438/438 [=====] - ETA: 0s - loss: 0.8216 - accuracy: 0.6948
Epoch 3: val_accuracy improved from 0.67641 to 0.75571, saving model to checkpoints_cnn_2\weights_epoch_03.h5
438/438 [=====] - 165s 376ms/step - loss: 0.8216 - accuracy: 0.6948 - val_loss: 0.6753 - val_accuracy: 0.7557
Epoch 4/20
438/438 [=====] - ETA: 0s - loss: 0.7598 - accuracy: 0.7258
Epoch 4: val_accuracy improved from 0.75571 to 0.77991, saving model to checkpoints_cnn_2\weights_epoch_04.h5
438/438 [=====] - 155s 354ms/step - loss: 0.7598 - accuracy: 0.7258 - val_loss: 0.6280 - val_accuracy: 0.7799
Epoch 5/20
438/438 [=====] - ETA: 0s - loss: 0.7020 - accuracy: 0.7500
Epoch 5: val_accuracy improved from 0.77991 to 0.78931, saving model to checkpoints_cnn_2\weights_epoch_05.h5
438/438 [=====] - 157s 357ms/step - loss: 0.7020 - accuracy: 0.7500 - val_loss: 0.6087 - val_accuracy: 0.7893
Epoch 6/20
438/438 [=====] - ETA: 0s - loss: 0.6572 - accuracy: 0.7659
Epoch 6: val_accuracy did not improve from 0.78931
438/438 [=====] - 159s 363ms/step - loss: 0.6572 - accuracy: 0.7659 - val_loss: 0.8201 - val_accuracy: 0.7151
Epoch 7/20
438/438 [=====] - ETA: 0s - loss: 0.6300 - accuracy: 0.7778
Epoch 7: val_accuracy improved from 0.78931 to 0.79301, saving model to checkpoints_cnn_2\weights_epoch_07.h5
438/438 [=====] - 157s 358ms/step - loss: 0.6300 - accuracy: 0.7778 - val_loss: 0.5815 - val_accuracy: 0.7930
Epoch 8/20
438/438 [=====] - ETA: 0s - loss: 0.6163 - accuracy: 0.7810
Epoch 8: val_accuracy improved from 0.79301 to 0.81183, saving model to checkpoints_cnn_2\weights_epoch_08.h5
438/438 [=====] - 156s 355ms/step - loss: 0.6163 - accuracy: 0.7810 - val_loss: 0.5401 - val_accuracy: 0.8118
Epoch 9/20
438/438 [=====] - ETA: 0s - loss: 0.5906 - accuracy: 0.7908
Epoch 9: val_accuracy improved from 0.81183 to 0.82124, saving model to checkpoints_cnn_2\weights_epoch_09.h5
438/438 [=====] - 156s 355ms/step - loss: 0.5906 - accuracy:
```

```
acy: 0.7908 - val_loss: 0.5113 - val_accuracy: 0.8212
Epoch 10/20
438/438 [=====] - ETA: 0s - loss: 0.5695 - accuracy: 0.7970
Epoch 10: val_accuracy did not improve from 0.82124
438/438 [=====] - 155s 354ms/step - loss: 0.5695 - accuracy: 0.7970 - val_loss: 0.5159 - val_accuracy: 0.8172
Epoch 11/20
438/438 [=====] - ETA: 0s - loss: 0.5509 - accuracy: 0.8081
Epoch 11: val_accuracy did not improve from 0.82124
438/438 [=====] - 157s 357ms/step - loss: 0.5509 - accuracy: 0.8081 - val_loss: 0.5159 - val_accuracy: 0.8196
Epoch 12/20
438/438 [=====] - ETA: 0s - loss: 0.5374 - accuracy: 0.8087
Epoch 12: val_accuracy did not improve from 0.82124
438/438 [=====] - 155s 354ms/step - loss: 0.5374 - accuracy: 0.8087 - val_loss: 0.5692 - val_accuracy: 0.8007
Epoch 13/20
438/438 [=====] - ETA: 0s - loss: 0.5241 - accuracy: 0.8164
Epoch 13: val_accuracy did not improve from 0.82124
438/438 [=====] - 156s 355ms/step - loss: 0.5241 - accuracy: 0.8164 - val_loss: 0.5641 - val_accuracy: 0.7997
Epoch 14/20
438/438 [=====] - ETA: 0s - loss: 0.5191 - accuracy: 0.8137
Epoch 14: val_accuracy did not improve from 0.82124
Restoring model weights from the end of the best epoch: 9.
438/438 [=====] - 155s 354ms/step - loss: 0.5191 - accuracy: 0.8137 - val_loss: 0.5950 - val_accuracy: 0.8004
Epoch 14: early stopping
```

```
In [16]: model_3 = build_cnn(nb_conv=3)

checkpoint, early_stop, ckpt_dir = get_callbacks(3)

history_3 = model_3.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // BATCH_SIZE,
    epochs=20,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // BATCH_SIZE,
    callbacks=[checkpoint, early_stop],
    verbose=1
)
```

Epoch 1/20
438/438 [=====] - ETA: 0s - loss: 1.1100 - accuracy: 0.5604
Epoch 1: val_accuracy improved from -inf to 0.68750, saving model to checkpoints_cnn_3\weights_epoch_01.h5
438/438 [=====] - 173s 393ms/step - loss: 1.1100 - accuracy: 0.5604 - val_loss: 0.8216 - val_accuracy: 0.6875
Epoch 2/20
438/438 [=====] - ETA: 0s - loss: 0.8551 - accuracy: 0.6808
Epoch 2: val_accuracy improved from 0.68750 to 0.72245, saving model to checkpoints_cnn_3\weights_epoch_02.h5
438/438 [=====] - 171s 389ms/step - loss: 0.8551 - accuracy: 0.6808 - val_loss: 0.7551 - val_accuracy: 0.7224
Epoch 3/20
438/438 [=====] - ETA: 0s - loss: 0.7460 - accuracy: 0.7348
Epoch 3: val_accuracy improved from 0.72245 to 0.79368, saving model to checkpoints_cnn_3\weights_epoch_03.h5
438/438 [=====] - 175s 399ms/step - loss: 0.7460 - accuracy: 0.7348 - val_loss: 0.5849 - val_accuracy: 0.7937
Epoch 4/20
438/438 [=====] - ETA: 0s - loss: 0.6552 - accuracy: 0.7700
Epoch 4: val_accuracy improved from 0.79368 to 0.80544, saving model to checkpoints_cnn_3\weights_epoch_04.h5
438/438 [=====] - 170s 389ms/step - loss: 0.6552 - accuracy: 0.7700 - val_loss: 0.5199 - val_accuracy: 0.8054
Epoch 5/20
438/438 [=====] - ETA: 0s - loss: 0.6082 - accuracy: 0.7862
Epoch 5: val_accuracy improved from 0.80544 to 0.83367, saving model to checkpoints_cnn_3\weights_epoch_05.h5
438/438 [=====] - 170s 389ms/step - loss: 0.6082 - accuracy: 0.7862 - val_loss: 0.4707 - val_accuracy: 0.8337
Epoch 6/20
438/438 [=====] - ETA: 0s - loss: 0.5526 - accuracy: 0.8070
Epoch 6: val_accuracy improved from 0.83367 to 0.84845, saving model to checkpoints_cnn_3\weights_epoch_06.h5
438/438 [=====] - 169s 386ms/step - loss: 0.5526 - accuracy: 0.8070 - val_loss: 0.4491 - val_accuracy: 0.8485
Epoch 7/20
438/438 [=====] - ETA: 0s - loss: 0.5237 - accuracy: 0.8144
Epoch 7: val_accuracy did not improve from 0.84845
438/438 [=====] - 169s 386ms/step - loss: 0.5237 - accuracy: 0.8144 - val_loss: 0.5078 - val_accuracy: 0.8212
Epoch 8/20
438/438 [=====] - ETA: 0s - loss: 0.4971 - accuracy: 0.8238
Epoch 8: val_accuracy did not improve from 0.84845
438/438 [=====] - 169s 386ms/step - loss: 0.4971 - accuracy: 0.8238 - val_loss: 0.5472 - val_accuracy: 0.8243
Epoch 9/20
438/438 [=====] - ETA: 0s - loss: 0.4860 - accuracy: 0.8326
Epoch 9: val_accuracy did not improve from 0.84845
438/438 [=====] - 169s 387ms/step - loss: 0.4860 - accuracy: 0.8326 - val_loss: 0.4484 - val_accuracy: 0.8441

```

Epoch 10/20
438/438 [=====] - ETA: 0s - loss: 0.4667 - accuracy: 0.8
398
Epoch 10: val_accuracy did not improve from 0.84845
438/438 [=====] - 169s 385ms/step - loss: 0.4667 - accuracy: 0.8398 - val_loss: 0.4371 - val_accuracy: 0.8458
Epoch 11/20
438/438 [=====] - ETA: 0s - loss: 0.4522 - accuracy: 0.8
425
Epoch 11: val_accuracy did not improve from 0.84845
Restoring model weights from the end of the best epoch: 6.
438/438 [=====] - 169s 386ms/step - loss: 0.4522 - accuracy: 0.8425 - val_loss: 0.4533 - val_accuracy: 0.8451
Epoch 11: early stopping

```

In [17]: `import matplotlib.pyplot as plt`

```

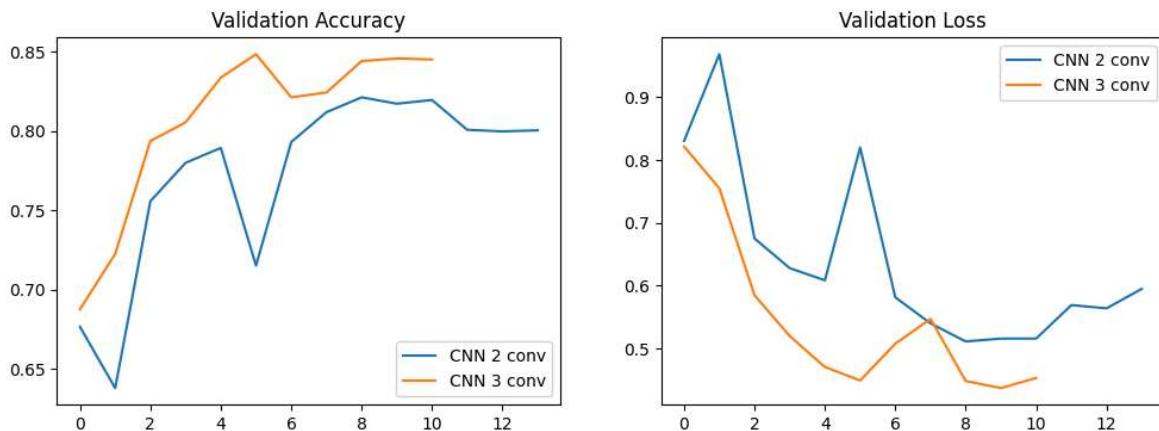
plt.figure(figsize=(12,4))

plt.subplot(1,2,1)
plt.plot(history_2.history["val_accuracy"], label="CNN 2 conv")
plt.plot(history_3.history["val_accuracy"], label="CNN 3 conv")
plt.title("Validation Accuracy")
plt.legend()

plt.subplot(1,2,2)
plt.plot(history_2.history["val_loss"], label="CNN 2 conv")
plt.plot(history_3.history["val_loss"], label="CNN 3 conv")
plt.title("Validation Loss")
plt.legend()

plt.show()

```



Comparaison CNN 2 convolutions vs CNN 3 convolutions

L'entraînement des deux architectures montre que l'ajout d'une troisième couche de convolution améliore significativement les performances sur le dataset Intel Image Classification.

- CNN 2 convolutions : atteint une précision maximale sur l'ensemble de validation d'environ 82,1 %, avec une perte de validation autour de 0,51.
- CNN 3 convolutions : atteint une précision maximale sur l'ensemble de validation d'environ 84,8 %, avec une perte de validation plus faible (~0,45).

Le graphique comparatif montre également que le modèle à 3 convolutions converge plus rapidement et généralise mieux, tout en maintenant un overfitting limité grâce à l'early stopping et au Dropout.

Donc avec de la troisième couche de convolution permet au CNN de capturer des caractéristiques plus complexes des images, ce qui améliore la précision globale et la stabilité du modèle. Pour ce dataset multiclass, le CNN à 3 convolutions est donc préférable.