

Si ModuleNotFoundError: No module named 'tqdm' ou 'GPUUtil' faire !pip install ... dans une cellule code

```
In [1]: import os
import numpy as np
import joblib
import cv2

import matplotlib.pyplot as plt
from tensorflow import keras
from sklearn.cluster import MiniBatchKMeans
from sklearn.svm import LinearSVC
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

import time
import GPUUtil
import psutil
```

```
In [2]: CHECKPOINT_DIR = "checkpoints_ORB"
os.makedirs(CHECKPOINT_DIR, exist_ok=True)
```

```
In [3]: UPSCALE_TO = 96
N_FEATURES = 500
VOCAB_SIZE = 300
KMEANS_BATCH = 4096
RANDOM_STATE = 42
```

```
In [4]: class PerformanceTracker:
    def __init__(self):
        self.start_time = None
        self.logs = {}

    def start(self, name):
        self.start_time = time.time()
        self.cpu_start = psutil.cpu_percent(interval=None)
        self.mem_start = psutil.virtual_memory().used

    def stop(self, name):
        elapsed = time.time() - self.start_time
        cpu = psutil.cpu_percent(interval=None)
        mem = psutil.virtual_memory().used
        self.logs[name] = {
            "time_sec": elapsed,
            "cpu_percent": cpu,
            "ram_mb": mem / 1024**2
        }
        print(f"[{name}] Time: {elapsed:.1f}s | CPU: {cpu}% | RAM: {mem/1024**2}:
```

```
In [5]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

print(x_train.shape, y_train.shape)
print(x_test.shape, y_test.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-dataset
s/mnist.npz
11490434/11490434 [=====] - 1s 0us/step
(60000, 28, 28) (60000,)
(10000, 28, 28) (10000,)
```

```
In [6]: x_train_u8 = x_train.astype(np.uint8)
x_test_u8 = x_test.astype(np.uint8)

def preprocess_for_orb(img_u8, upscale_to=96):
    if upscale_to is not None and upscale_to != img_u8.shape[0]:
        img_u8 = cv2.resize(img_u8, (upscale_to, upscale_to), interpolation=cv2.INTER_LINEAR)
    return img_u8
orb = cv2.ORB_create(nfeatures=N_FEATURES)
```

```
In [7]: def extract_orb_desc_from_array(images_u8, upscale_to=96, max_images=None):
    desc_list = []
    valid_idx = []

    n = len(images_u8) if max_images is None else min(len(images_u8), max_images)

    for i in range(n):
        img = preprocess_for_orb(images_u8[i], upscale_to=upscale_to)
        kp, desc = orb.detectAndCompute(img, None)
        if desc is None or len(desc) == 0:
            continue
        desc_list.append(desc) # (n_kp, 32)
        valid_idx.append(i)

    return desc_list, np.array(valid_idx, dtype=np.int32)
```

```
In [8]: tracker = PerformanceTracker()

tracker.start("ORB extraction (train+test)")
train_desc_list, train_valid_idx = extract_orb_desc_from_array(x_train_u8, upscale_to=96)
test_desc_list, test_valid_idx = extract_orb_desc_from_array(x_test_u8, upscale_to=96)
tracker.stop("ORB extraction (train+test)")
y_train_f = y_train[train_valid_idx]
y_test_f = y_test[test_valid_idx]
```

[ORB extraction (train+test)] Time: 33.8s | CPU: 20.6% | RAM: 7374 MB

```
In [9]: all_train_desc = np.vstack(train_desc_list).astype(np.float32) # (N_total_kp, 32)
print("Total train descriptors:", all_train_desc.shape)

kmeans = MiniBatchKMeans(
    n_clusters=VOCAB_SIZE,
    batch_size=KMEANS_BATCH,
    random_state=42,
    n_init="auto"
)

tracker.start("KMeans vocab")
kmeans.fit(all_train_desc)
tracker.stop("KMeans vocab")

joblib.dump(kmeans, os.path.join(CHECKPOINT_DIR, "kmeans_vocab_mnist.pkl"))
print("Saved vocab to:", os.path.join(CHECKPOINT_DIR, "kmeans_vocab_mnist.pkl"))
```

Total train descriptors: (1024077, 32)
 [KMeans vocab] Time: 4.7s | CPU: 43.4% | RAM: 7075 MB
 Saved vocab to: checkpoints_ORB\kmeans_vocab_mnist.pkl

```
In [10]: def bovw_histograms(desc_list, kmeans, vocab_size):
    X = np.zeros((len(desc_list), vocab_size), dtype=np.float32)

    for i, desc in enumerate(desc_list):
        words = kmeans.predict(desc.astype(np.float32))
        hist = np.bincount(words, minlength=vocab_size).astype(np.float32)

        # Normalisation L2 (souvent mieux pour SVM)
        norm = np.linalg.norm(hist) + 1e-12
        X[i] = hist / norm

    return X

tracker.start("BoVW histograms")
X_train = bovw_histograms(train_desc_list, kmeans, VOCAB_SIZE)
X_test = bovw_histograms(test_desc_list, kmeans, VOCAB_SIZE)
tracker.stop("BoVW histograms")
print("X_train:", X_train.shape)
print("X_test :", X_test.shape)
```

[BoVW histograms] Time: 12.7s | CPU: 23.3% | RAM: 7153 MB
 X_train: (59815, 300)
 X_test : (9971, 300)

```
In [11]: svm = LinearSVC(dual="auto", random_state=42)
tracker.start("SVM training")
svm.fit(X_train, y_train_f)
tracker.stop("SVM training")

joblib.dump(svm, os.path.join(CHECKPOINT_DIR, "svm_bovw_orb_mnist.pkl"))
print("Saved SVM to:", os.path.join(CHECKPOINT_DIR, "svm_bovw_orb_mnist.pkl"))
```

[SVM training] Time: 3.5s | CPU: 24.0% | RAM: 7166 MB
 Saved SVM to: checkpoints_ORB\svm_bovw_orb_mnist.pkl

```
In [12]: y_pred = svm.predict(X_test)

acc = accuracy_score(y_test_f, y_pred)
print("Accuracy:", acc)

print("\nClassification report:")
print(classification_report(y_test_f, y_pred))

print("\nConfusion matrix:")
print(confusion_matrix(y_test_f, y_pred))
```

Accuracy: 0.6323337679269883

Classification report:

	precision	recall	f1-score	support
0	0.63	0.74	0.68	966
1	0.78	0.95	0.86	1121
2	0.54	0.43	0.48	1031
3	0.65	0.72	0.68	1010
4	0.58	0.52	0.54	982
5	0.55	0.59	0.57	892
6	0.62	0.59	0.61	958
7	0.64	0.49	0.55	1028
8	0.65	0.69	0.67	974
9	0.61	0.57	0.59	1009
accuracy			0.63	9971
macro avg	0.62	0.63	0.62	9971
weighted avg	0.63	0.63	0.63	9971

Confusion matrix:

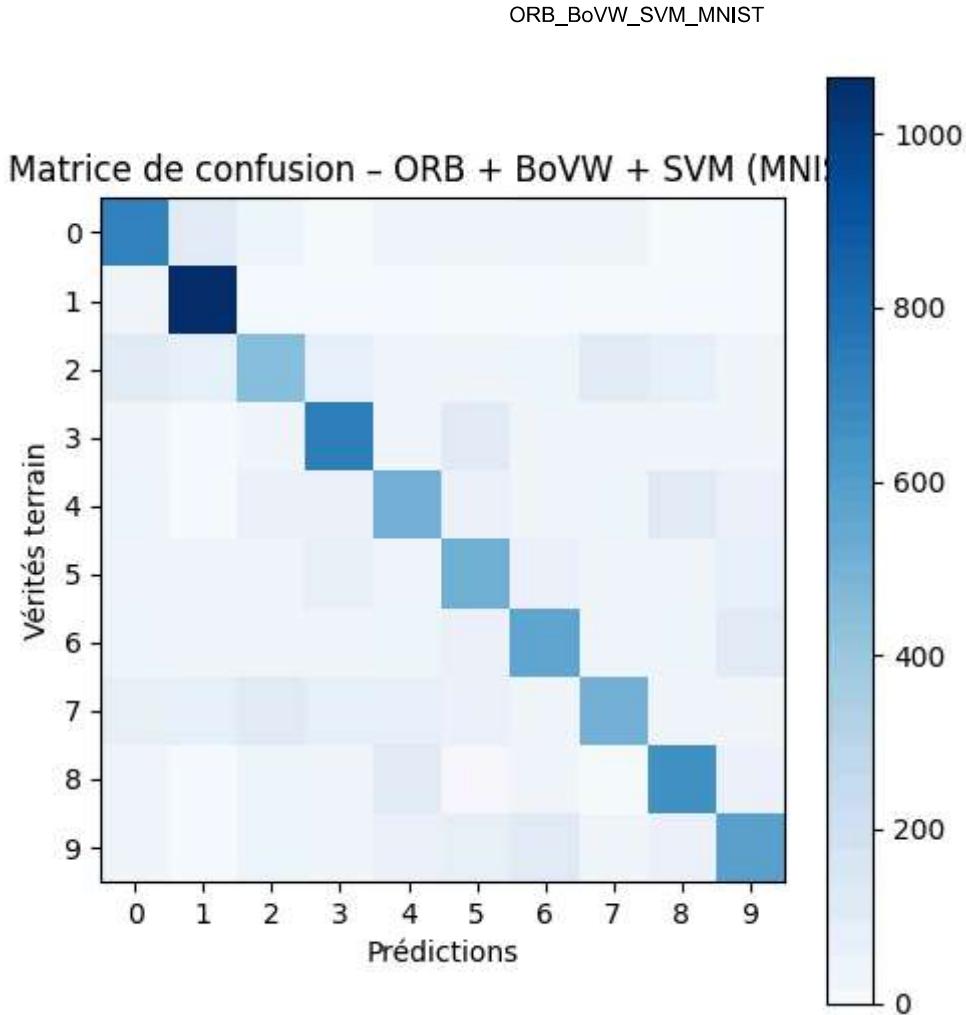
```
[[ 716 105 42   8  16 14  21  34   7   3]
 [ 34 1065  3   4   0   6   0   7   1   1]
 [ 113  67 447  79  32 35  46 111  78  23]
 [ 27   3  33 732  29 96  19  21  25  25]
 [ 40   3  58  64 507 57  37  43 112  61]
 [ 20  15  33  90  30 524  56  27  22  75]
 [ 41  15  36  26  39 63  564  24  45 105]
 [ 91  81  97  71  67 65  13  506  24  13]
 [ 35   0  44  24 101 10  30   6 669  55]
 [ 28   3  38  35  60 85 117  17  51 575]]
```

```
In [13]: cm = confusion_matrix(y_test_f, y_pred)

plt.figure(figsize=(5,5))
plt.imshow(cm, cmap="Blues")
plt.title("Matrice de confusion - ORB + BoVW + SVM (MNIST)")
plt.xlabel("Prédictions")
plt.ylabel("Vérités terrain")
plt.colorbar()

# MNIST : classes = chiffres 0 à 9
classes = list(range(10))
plt.xticks(classes, classes)
plt.yticks(classes, classes)

plt.tight_layout()
plt.show()
```



Conclusion

L'application de la méthode ORB combinée au Bag of Visual Words et à un SVM sur le jeu de données MNIST a permis de tester une approche générique sur un problème de reconnaissance de chiffres manuscrits. Toutefois, ce jeu de données, composé d'images simples, centrées et peu texturées, se prête mal à l'utilisation de descripteurs locaux comme ORB.

Les points clés détectés sont peu nombreux et parfois peu informatifs, ce qui limite fortement la qualité de la représentation BoVW. Par ailleurs, le coût de calcul lié à l'extraction des descripteurs et au clustering reste non négligeable sur CPU, sans apporter de réel avantage par rapport à des méthodes plus simples.

Dans ce contexte, cette approche apparaît surdimensionnée et inefficace pour MNIST.

In []: